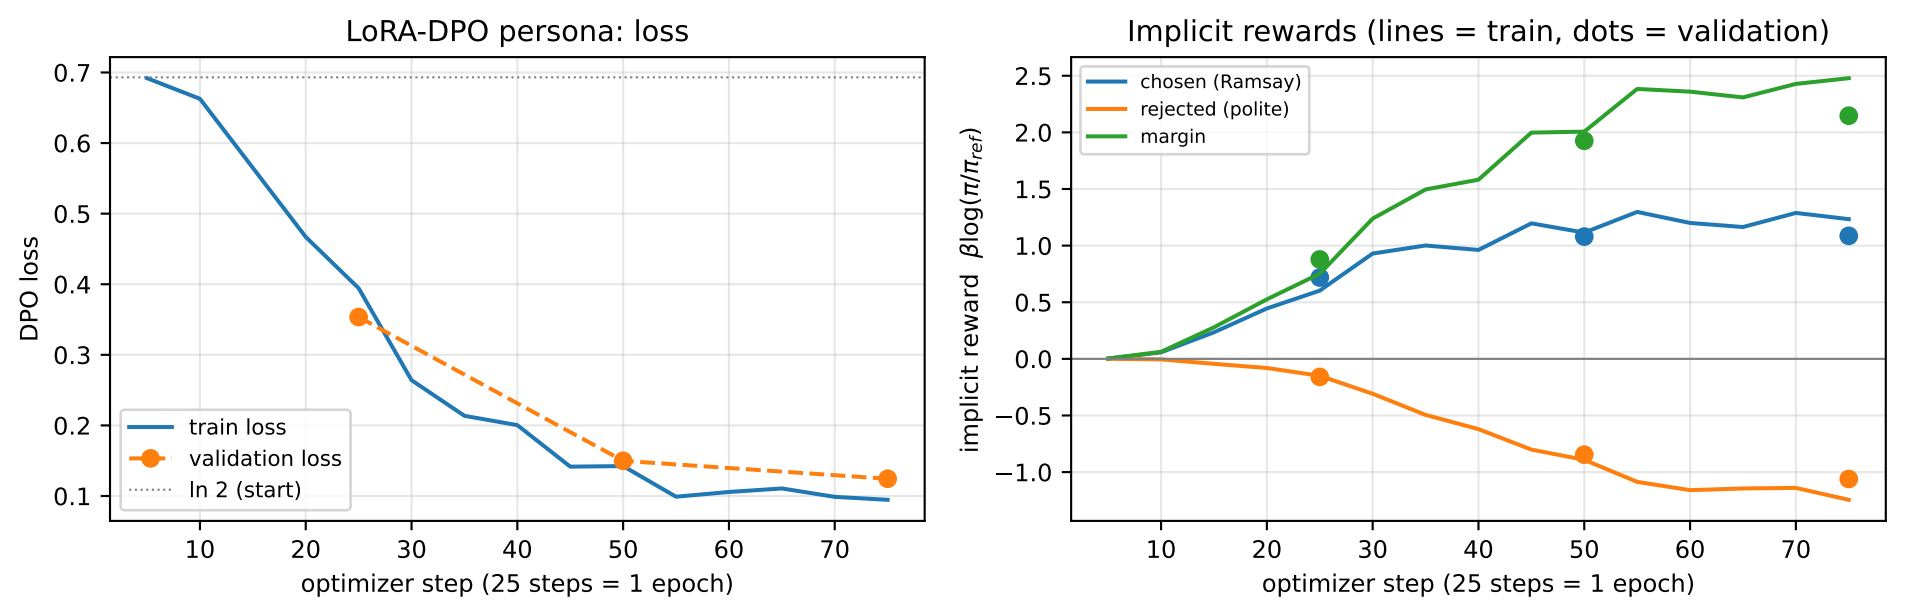
LoRA-DPO persona: loss
DPO loss
optimizer step (25 steps = 1 epoch)
train loss
validation loss
ln 2 (start)
Implicit rewards (lines = train, dots = validation)
implicit reward $\beta \log(\pi/\pi_{ref})$
optimizer step (25 steps = 1 epoch)
chosen (Ramsay)
rejected (polite)
margin